

```
[35]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
import seaborn as sns
import matplotlib.pyplot as plt
```

```
[39]: # Specify the file path
file = "D:\\intership\\New folder\\car prediction\\CarPrice_Assignment.csv"

data = pd.read_csv(file)

data.head()
```

```
[39]:
```

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	engineLocation	wheelbase	...	enginesize	fuelsystem	boreratio	stroke	cc
0	1	3	alfa-romero giulia	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	2.68	
1	2	3	alfa-romero stelvio	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	2.68	
2	3	1	alfa-romero Quadrifoglio	gas	std	two	hatchback	rwd	front	94.5	...	152	mpfi	2.68	3.47	
3	4	2	audi 100 ls	gas	std	four	sedan	fwd	front	99.8	...	109	mpfi	3.19	3.40	
4	5	2	audi 100ls	gas	std	four	sedan	4wd	front	99.4	...	136	mpfi	3.19	3.40	

5 rows × 26 columns

```
[40]: # Check missing values
data.isnull().sum()
```

```
[40]: car_ID      0
symboling    0
CarName      0
fueltype     0
aspiration   0
doornumber   0
carbody      0
drivewheel   0
engineLocation 0
wheelbase    0
carlength    0
carwidth     0
carheight    0
curbweight   0
engineType   0
cylindernumber 0
enginesize   0
fuelsystem   0
boreratio    0
stroke       0
compressionratio 0
horsepower   0
peakrpm      0
citympg      0
highwaympg   0
price        0
dtype: int64
```

```
*[42]: # Handle missing values
data = data.dropna()
data
```

```
[42]:
```

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	engineLocation	wheelbase	...	enginesize	fuelsystem	boreratio	stroke	cc
0	1	3	alfa-romero giulia	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	2.68	
1	2	3	alfa-romero stelvio	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	2.68	
2	3	1	alfa-romero Quadrifoglio	gas	std	two	hatchback	rwd	front	94.5	...	152	mpfi	2.68	3.47	
3	4	2	audi 100 ls	gas	std	four	sedan	fwd	front	99.8	...	109	mpfi	3.19	3.40	
4	5	2	audi 100ls	gas	std	four	sedan	4wd	front	99.4	...	136	mpfi	3.19	3.40	
...
200	201	-1	volvo 145e (sw)	gas	std	four	sedan	rwd	front	109.1	...	141	mpfi	3.78	3.15	
201	202	-1	volvo 144ea	gas	turbo	four	sedan	rwd	front	109.1	...	141	mpfi	3.78	3.15	
202	203	-1	volvo 244dl	gas	std	four	sedan	rwd	front	109.1	...	173	mpfi	3.58	2.87	
203	204	-1	volvo 246	diesel	turbo	four	sedan	rwd	front	109.1	...	145	idi	3.01	3.40	
204	205	-1	volvo 264gl	gas	turbo	four	sedan	rwd	front	109.1	...	141	mpfi	3.78	3.15	

205 rows × 26 columns

```
[49]: # Convert categorical variables to numerical values using one-hot encoding
data = pd.get_dummies(data, drop_first=True)

data.columns
```

```
[49]: Index(['car_ID', 'symboling', 'wheelbase', 'carlength', 'carwidth',
```

```

        'carheight', 'curbweight', 'enginesize', 'boreratio', 'stroke',
        ...
        'cylindernumber_three', 'cylindernumber_twelve', 'cylindernumber_two',
        'fuelsystem_2bbl', 'fuelsystem_4bbl', 'fuelsystem_idi',
        'fuelsystem_mfi', 'fuelsystem_mphi', 'fuelsystem_spdi',
        'fuelsystem_spfi'],
        dtype='object', length=191)
[50]: data
[50]:
   car_ID  symboling  wheelbase  carlength  carwidth  carheight  curbweight  enginesize  boreratio  stroke  ...  cylindernumber_three  cylindernumber_twelve  cylindern
0        1         3         88.6        168.8         64.1         48.8        2548         130         3.47         2.68  ...                False                False
1        2         3         88.6        168.8         64.1         48.8        2548         130         3.47         2.68  ...                False                False
2        3         1         94.5        171.2         65.5         52.4        2823         152         2.68         3.47  ...                False                False
3        4         2         99.8        176.6         66.2         54.3        2337         109         3.19         3.40  ...                False                False
4        5         2         99.4        176.6         66.4         54.3        2824         136         3.19         3.40  ...                False                False
...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...  ...      ...      ...
200     201        -1        109.1        188.8         68.9         55.5        2952         141         3.78         3.15  ...                False                False
201     202        -1        109.1        188.8         68.8         55.5        3049         141         3.78         3.15  ...                False                False
202     203        -1        109.1        188.8         68.9         55.5        3012         173         3.58         2.87  ...                False                False
203     204        -1        109.1        188.8         68.9         55.5        3217         145         3.01         3.40  ...                False                False
204     205        -1        109.1        188.8         68.9         55.5        3062         141         3.78         3.15  ...                False                False

```

205 rows × 191 columns

```

[52]: # Define features and target variable
X = data.drop(columns=['price'])
y = data['price']

# Display the features and target variable
X

```

```

[52]: 0      13495.0
      1      16500.0
      2      16500.0
      3      13950.0
      4      17450.0
      ...
     200     16845.0
     201     19045.0
     202     21485.0
     203     22470.0
     204     22625.0
      Name: price, Length: 205, dtype: float64

```

```

[53]: y
[53]: 0      13495.0
      1      16500.0
      2      16500.0
      3      13950.0
      4      17450.0
      ...
     200     16845.0
     201     19045.0
     202     21485.0
     203     22470.0
     204     22625.0
      Name: price, Length: 205, dtype: float64

```

```

[54]: # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

```

```

[56]: # Initialize the Linear Regression model
model = LinearRegression()

# Train the model
model.fit(X_train, y_train)

```

```

[56]: • LinearRegression
      LinearRegression()

```

```


[58]: # Make predictions on the test set
y_pred = ln_model.predict(X_test)


# Calculate the Mean Squared Error (MSE) and R^2 score
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
print(f'Mean Squared Error: {mse:.2f}')
print(f'R^2 Score: {r2}')

# Plot the actual vs predicted prices
plt.scatter(y_test, y_pred)
plt.xlabel('Actual Prices')
plt.ylabel('Predicted Prices')
plt.title('Actual vs Predicted Prices')
plt.show()

```

Mean Squared Error: 196438429.48
R^2 Score: -1.8352484848727344





[]:

[]:

